



ANSIBLE

Authored and Presented by: Amol Shende

Agenda

- Prerequisite- ssh, configuration management
- Need of Ansible
- What is Ansible?
- Architecture
- Key Concepts
- Features
- Cons
- Some Advance Concepts
- Case Studies

SSH (Secure Shell)

- SSH is cryptographic network protocol
- Used for secure data-communication, remote login, remote command execution and other secure network services
- Standard port 22
- Ansible uses SSH for performing its tasks
- Authentication is done using username-password or public-key cryptography algorithm
- Utilities- SCP, SFTP, x-11 forwarding

Configuration Management

- What so you need to build good software? (**System, good environment**)
- It is process to setup such system and environment in well documented manner
- when applied to system, provides visibility, performance , maintainability and reduce cost and risk
- It makes sure that system is behaving as intended and documented way to support project life cycle
- Two ways to do 1) Manual 2) Automatic


Need Of Ansible (Manual Vs Automatic)

Shell Scripting

- 1) Long and Tedious scripts
- 2) Difficult to maintain
- 3) Run scripts manually one machine-one script at a time
- 4) Need lots of time
- 5) Difficult to track
- 6) Can not handle Hybrid Infrastructure
- 7) Poor Documentation
- 8) Need Linux expert
- 9) No Idempotent

Ansible

- 1) Short and readable
- 2) Easy to maintain
- 3) Can run multiple scripts on 1000s of nodes at time
- 4) Very less time
- 5) Easy to track
- 6) Easy way to handle Hybrid Infrastructure
- 7) Proper and structured documentation
- 8) Basic Linux knowledge is sufficient
- 9) Idempotent



What is Ansible?

- Configuration Management Tool
- Developed by **Michael DeHaan**
- Initial Release February 2012

History

- 1st generation
 - CF Engine
- 2nd generation
 - Puppet
 - Chef
- 3rd generation
 - Ansible
 - Salt Stack



Who Uses?

Who uses ANSIBLE?

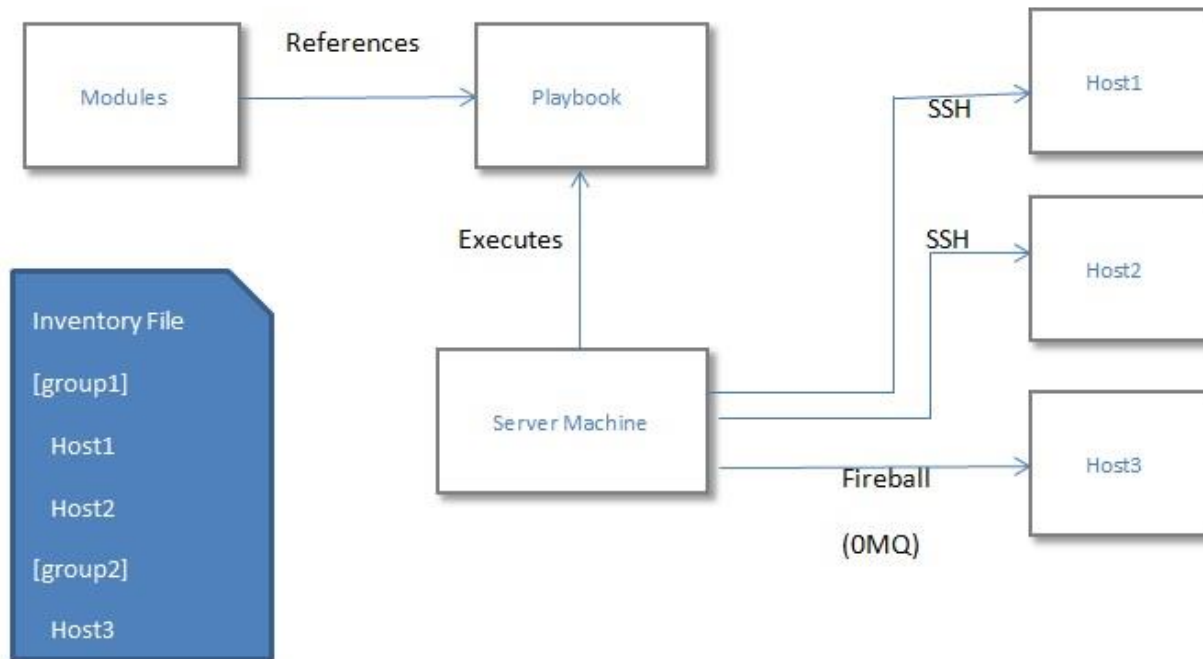


Goals

- Configuration Management
- Application Deployment
- Cloud Provisioning
- AD-HOC Task execution
- Continuous Deployment
- zero downtime rolling updates



Architecture



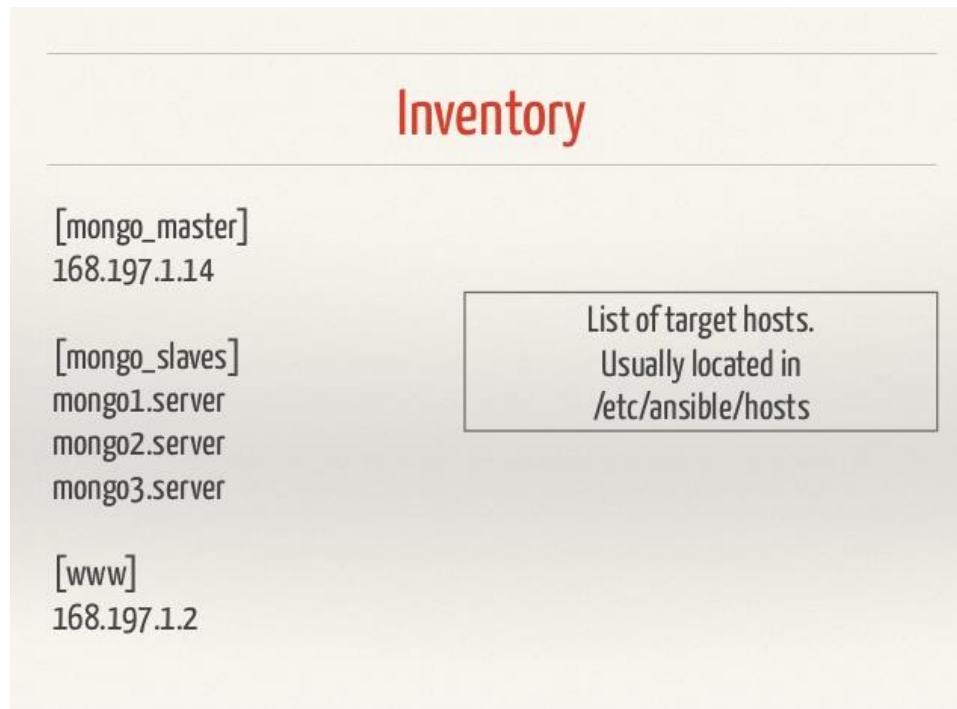
Key Concepts

- Inventory
- Modules
- Ad Hoc Commands
- Playbooks
 - Roles
 - Tasks
 - Variables
 - Templates
 - Handlers



Inventory

- Define a list of target hosts
- usually located in `/etc/ansible/hosts`



Inventory

- If remote host ssh is not running on default port(22) then,

hostname : ssh port

- Using alias for hostname

Alias_name ansible_ssh_port=5326 ansible_ssh_host=172.27.26.9

- Connection type and user on a per host basis:

localhost ansible_connection=local

other1.example.com ansible_connection=ssh ansible_ssh_user=amol

- Group Variables : can assign variable to entire group of hosts

[atlanta]

host1

host2

[atlanta:vars]

ntp_server=ntp.atlanta.example.com

proxy=proxy.atlanta.example.com

Note: Best practice is not to keep group and host vars in inventory file. Instead keep in /etc/ansible/group_vars/file_name folder.

➤ Dynamic Inventory

Ex. local inventories

Ansible-playbook -i /inventory/file/path playbook.yml

Ex. cloud inventories (AWS EC2)

Ansible-playbook -i /ec2.py/file/path playbook.yml

Note: 1) To make API call AWS, you need to configure Boto(Python Interface for AWS)

2) Also need to export 2 environment variables

export AWS_ACCESS_KEY_ID='AK123'

export AWS_SECRET_ACCESS_KEY_ID='ab123'

➤ Hybrid Infrastructure Inventory

Ansible-playbook -i /path/of/inventory_file_directory playbook.yml


Host Defining Patterns

- All / *
- 192.168.1.*
- webserver:dbserver
- webserver:!phoenix
- webserver:&staging
- webserver[0]
- webserver[0:25]
- ansible-playbook site.yml --limit @retry_hosts.txt

Modules

- They are like libraries
- Playbook internally calls these modules
- All core modules written in Python
- can be written in any language as long as they output JSON
- take parameters and conditions to define desired state
- handles processing of system resources, services, packages, files, etc. in idempotent fashion
- ansible comes preloaded with a lots of modules

```
- name: Install apache  
  apt: pkg=apache2 state=present
```



Ad-Hoc Commands

- These are ansible simple commands for simple tasks, such as making sure a service is running, or to trigger updates and reboots.
- run a single, one-off command
- run on a full or partial inventory
- no need to save for later
- Format : `ansible <host pattern> <module> <arguments>`



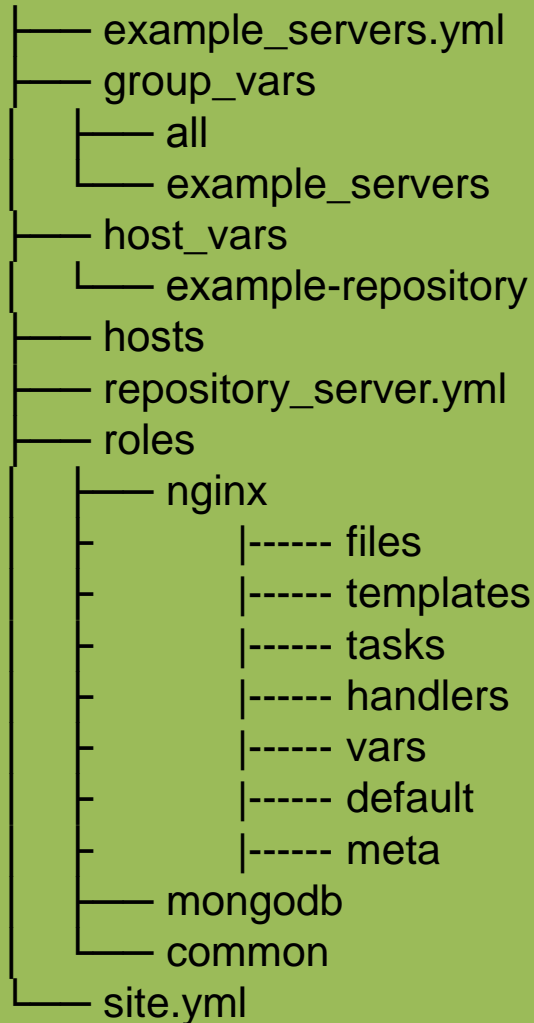
Ex. `ansible all -m user -a "name=joe password=wat"`

Playbook

- It is collection of plays (tasks)
- For more powerful configuration management
- Kept in source control, developed, validated
- Declare configurations of more complex multi-system environments
- Arrange and run tasks synchronously or asynchronously



Anatomy Of Playbook



Play Example

```
---  
- hosts: mongo_masters  
  user: root  
  tasks:  
    - name: Install apache2  
      apt: name=apache2 state=present  
      notify:  
        - restart apache  
    - name: ensure apache is running  
      service: name=apache2 state=started  
  handlers:  
    - name: restart apache  
      service: name=apache2 state=restarted
```

Running Playbook

ansible-playbook test.yml

ansible-playbook test.yml --ask-pass --ask-sudo-pass

ansible-playbook -i hosts -k -K test.yml

Tasks

- Each playbook contains a list of tasks.
- Tasks are executed in order, one at a time, against all machines matched by the host pattern, before moving on to the next task
- These tasks calls ansible modules to perform task.

Variables

- Vars can be defined at many levels (default, role, playbook, inventory)
- Then you can reference them-

1) on command line

```
$ ansible-playbook site.yml --extra-vars="username=example domain=example.org"
```

2) in a task

```
- name: Create Vhost User  
  user: name={{ username }} state=present
```

Variables cont..

3) in a template

```
server {  
    listen 80;  
    server_name www.{{ domain }};  
    root /home/{{ username }}/public_html;  
    index index.html index.php;  
    access_log /home/{{ username }}/logs/access.log;  
    error_log /home/{{ username }}/logs/error.log warn;  
    .....  
}
```

4) Complex Variables

Variables cont..

5) Magical Variables

- reserved and used to access info on other hosts

Ex. 1) hostvars - lets you ask about the variables of another host, including facts

```
{{ hostvars['test.example.com']['ansible_distribution'] }}
```

2) groups - list of all the groups (and hosts) in the inventory

```
{% for host in groups['app_servers'] %}
```

```
# something that applies to all app servers.
```

```
{% endfor %}
```

3) group_names - list (array) of all the groups the current host is in


6) Variables in other files

vars_files:

- /vars/external_vars.yml

Templates

- Ansible uses jinja2 templating format
- Templates are interpreted by jinja2 template engine
- fill variables in differently depending on conditions
- It allows Ansible to use some programming language features
- Powerful conditionals
- Loops and iterators



Handlers

- The things listed in the 'notify' section of a task are called handlers.
- Handlers are how you can perform various post-deployment tasks like restating some services when some of its configuration files changed. Handlers only get triggered if the state of the line actually changed
- These 'notify' actions are triggered at the end of each block of tasks in a playbook, and will only be triggered once even if notified by multiple different tasks.
- Notify handlers are always run in the order written.

Roles

- A feature of Ansible for encouraging reuse of code and best practices.
- Encapsulate a reusable service definition including Actions, Variables, Templates, Files
- Defined in standard directory structure (subproject-like)
- A host / group can have multiple roles
- A role can be applied to multiple groups
- You can upload or download roles from Ansible-Galaxy

```
▼ ansible-tutorial-roles
  ▼ roles
    ▼ apache
      ▼ files
        ► secure
      ▼ tasks
        ► main.yml
      ▼ templates
        ► service-ssl.j2
      ▼ vars
        ► main.yml
    ► common
    ► jdk7
    ► postgres
    ► service_base
    ► tomcat
    ► pizzamatic-roles.playbook
```

Roles

➤ Parameterized roles

roles:

```
- { role: foo_app_instance, dir: '/opt/a', port: 5000 }
```

➤ pre_tasks and post_tasks

Features

- 1) Agent less Architecture
- 2) Python based
- 3) YAML Syntax
- 4) Modules can be written in any language
- 5) Push based
- 6) Supports pull mode (git required)
- 7) Fireball mode/Parallel Mode/Accelerated mode
- 8) Encryption and security built in
- 9) Idempotent

Cons

- GUI is not well developed
- Came in 2012 , so not used by many
- Very less support for Windows

Ansible Window Support

- Starting in version 1.7, Ansible also contains support for managing Windows machines. This uses native powershell remoting, rather than SSH.
- In order for Ansible to manage your windows machines, you will have to enable Powershell remoting configured.
- A Linux control machine will be required to manage Windows hosts

Check Mode

- When ansible-playbook is executed with "- -check" it will not make any changes on remote systems. Instead will report what changes they would have made rather than making them
- "- - diff" flag to find out changes
- "-- list-host" flag to list affecting hosts
- "-- verbose" flag to see detailed output

Prompts

When running a playbook, you may wish to prompt the user for certain input, and can do so with the 'vars_prompt' section.

Ex.

```
- hosts: all
  remote_user: root
  vars:
    from: "camelot"
  vars_prompt:
    name: "what is your name?"
    quest: "what is your quest?"
    favcolor: "what is your favorite color?"
```

Tags

If you have a large playbook it may become useful to be able to run a specific part of the configuration without running the whole playbook.

tasks:

- name: install pkgs
yum: name={{ item }} state=installed

with_items:

- httpd
- memcached

tags:

- packages

```
aansible-playbook example.yml --tags "packages"  
ansible-playbook example.yml --skip-tags "packages"
```

Fireball Mode

- 1) Deprecated
- 2) Used to speed up ansible working.
- 3) Alternatives are Pipelining and Accelerated mode.

Accelerared Mode

- 1) Use only if you are unable to use pipelining feature of ansible.
- 2) Accelerated mode can be anywhere from 2-6x faster than SSH.
- 3) Launches Demon Process Over SSH
- 4) Uses Socket connection
- 5) `accelerate_port = 5099`, ---- default port
- 6) `accelerate_timeout = 30`, ----- timeout for receiving data from a client.(should be greater than 15s)
- 7) `accelerate_connect_timeout = 1.0`, ----- controls the timeout for the socket connect call. If 3 attempts fails then it will connect using default ssh or paramiko.

Pipelining

- 1) Enabling pipelining reduces the number of SSH operations required to execute a module on the remote server, by executing many ansible modules without actual file transfer.
- 2) This can result in a very significant performance improvement when enabled.
- 3) pipelining is better than accelerate mode for nearly all use cases.

Asynchronous Mode

- 1) By default, when you run playbook , the ssh connection remains open until the task is done on each node.
- 2) This may not always be desirable because you are blocking ssh connection or you may be running operations that take longer than the SSH timeout.
- 3) To resolve this problem Asynchronous Mode is used.
- 4) Here we kick off all modules at once and then poll until they are done.

Vault

“Vault” is a feature of ansible that allows encryption of ansible files and data

```
ansible-vault create foo.yml
ansible-vault edit foo.yml
ansible-vault rekey foo.yml bar.yml baz.yml
ansible-vault encrypt foo.yml bar.yml baz.yml
ansible-vault decrypt foo.yml bar.yml baz.yml
ansible-vault view foo.yml bar.yml baz.yml
ansible-playbook site.yml --ask-vault-pass
ansible-playbook site.yml --vault-password-file ~/.vault_pass.txt
```

Start and Step

➤ Start-at-task

Ex. `ansible-playbook playbook.yml --start-at-task="install packages"`

➤ Step

Ex. `ansible-playbook playbook.yml -step`

Perform task: install packages (y/n/c):

Task Delegation

1) Doing things on one host on behalf of another

Ex. - - hosts: webservers

tasks:

-Name: Add to load-balancer pool

command: add_to_loadbalancer_pool

delegate_to: 172.27.59.17

1) Suppose you running a playbook for group of 10 hosts. But you want particular task in playbook should run only once on particular host ,then you should use delegate_to option.

Ex. - task:

-Name: install nginx

apt: name=nginx state=present

run_once: true

delegate_to: 172.27.59.17

Any Questions?





Thank you!